

02_planar_data_classification_with_one_hidden_layer

April 7, 2025

1 Planar data classification with one hidden layer

In this exercise, we will build our first neural network which will have one hidden layer. We'll notice a big difference between this model and the one we implemented previously using logistic regression.

By the end of this assignment, you'll be able to:

- Implement a 2-class classification neural network with a single hidden layer
- Use units with a non-linear activation function, such as tanh
- Compute the cross entropy loss
- Implement forward and backward propagation

1.1 Table of Contents

- 1 - Packages
- 2 - Load the Dataset
 - Exercise 1
- 3 - Simple Logistic Regression
- 4 - Neural Network model
 - 4.1 - Defining the neural network structure
 - * Exercise 2 - layer_sizes
 - 4.2 - Initialize the model's parameters
 - * Exercise 3 - initialize_parameters
 - 4.3 - The Loop
 - * Exercise 4 - forward_propagation
 - 4.4 - Compute the Cost
 - * Exercise 5 - compute_cost
 - 4.5 - Implement Backpropagation
 - * Exercise 6 - backward_propagation
 - 4.6 - Update Parameters
 - * Exercise 7 - update_parameters
 - 4.7 - Integration
 - * Exercise 8 - nn_model
- 5 - Test the Model
 - 5.1 - Predict
 - * Exercise 9 - predict
 - 5.2 - Test the Model on the Planar Dataset
- 6 - Tuning hidden layer size (optional/ungraded exercise)

- 7- Performance on other datasets

1 - Packages

First import all the packages that we will need during this assignment.

- `numpy` is the fundamental package for scientific computing with Python.
- `sklearn` provides simple and efficient tools for data mining and data analysis.
- `matplotlib` is a library for plotting graphs in Python.
- `testCases` provides some test examples to assess the correctness of your functions
- `planar_utils` provide various useful functions used in this assignment

```
[1]: # Package imports
import numpy as np
import copy
import matplotlib.pyplot as plt
from testCases_v2 import *
from public_tests import *
import sklearn
import sklearn.datasets
import sklearn.linear_model
from planar_utils import plot_decision_boundary, sigmoid, load_planar_dataset, load_extra_datasets

%matplotlib inline


%load_ext autoreload
%autoreload 2
```

2 - Load the Dataset

```
[2]: X, Y = load_planar_dataset()
```

Visualize the dataset using `matplotlib`. The data looks like a “flower” with some red (label $y=0$) and some blue ($y=1$) points. Our goal is to build a model to fit this data. In other words, we want the classifier to define regions as either red or blue.

```
[3]: # Visualize the data:
plt.scatter(X[0, :], X[1, :], c=Y, s=40, cmap=plt.cm.Spectral);
```



We have the following:

- a numpy-array (matrix) `X` that contains the features (`x1`, `x2`)
- a numpy-array (vector) `Y` that contains the labels (`red:0`, `blue:1`).

First, let's get a better sense of what our data is like.

Exercise 1

How many training examples do we have? In addition, what is the **shape** of the variables `X` and `Y`?

Hint: How do we get the shape of a numpy array? ([help](#))

[4]: *# The structure of your code in this cell should be as follows:*

```
# shape_X = ...
# shape_Y = ...
# training set size
# m = ...

# YOUR CODE STARTS HERE

shape_X = X.shape      # X is of shape (num_px * num_px * 3, m)
shape_Y = Y.shape      # Y is of shape (1, m)
```

```

m = X.shape[1]          # Number of training examples is the number of columns
↪ in X

# YOUR CODE ENDS HERE

print ('The shape of X is: ' + str(shape_X))
print ('The shape of Y is: ' + str(shape_Y))
print ('I have m = %d training examples!' % (m))

```

The shape of X is: (2, 400)
The shape of Y is: (1, 400)
I have m = 400 training examples!

Expected Output:

shape of X

(2, 400)

shape of Y

(1, 400)

m

400

3 - Simple Logistic Regression

Before building a full neural network, let's check how logistic regression performs on this problem. We will use sklearn's built-in functions for this. Run the code below to train a logistic regression classifier on the dataset.

```

[5]: # Train the logistic regression classifier
     clf = sklearn.linear_model.LogisticRegressionCV();
     clf.fit(X.T, Y.T);

```

```

/home/codespace/.local/lib/python3.12/site-
packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-
vector y was passed when a 1d array was expected. Please change the shape of y
to (n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)

```

Let's take a look at the decision boundary of this model. Run the code below.

```

[6]: # Plot the decision boundary for logistic regression
     plot_decision_boundary(lambda x: clf.predict(x), X, Y)
     plt.title("Logistic Regression")

     # Print accuracy


```

```
LR_predictions = clf.predict(X.T)
print ('Accuracy of logistic regression: %d ' % float((np.dot(Y,LR_predictions) +
↪ np.dot(1-Y,1-LR_predictions))/float(Y.size)*100) +
      '% ' + "(percentage of correctly labelled datapoints)")
```

/tmp/ipykernel_13494/4242423965.py:7: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

```
print ('Accuracy of logistic regression: %d ' %
float((np.dot(Y,LR_predictions) +
np.dot(1-Y,1-LR_predictions))/float(Y.size)*100) +
```

Accuracy of logistic regression: 47 % (percentage of correctly labelled datapoints)



Expected Output:

Accuracy

47%

Interpretation: The dataset is not linearly separable, so logistic regression doesn't perform well. Hopefully a neural network will do better...

4 - Neural Network model

Logistic regression didn't work well on the flower dataset. So, we're going to train a Neural Network with a single hidden layer and see how that handles the same problem.

The model:

Mathematically:

For one example $x^{(i)}$:

$$z^{[1](i)} = W^{[1]}x^{(i)} + b^{[1]} \quad (1)$$

$$a^{[1](i)} = \tanh(z^{[1](i)}) \quad (2)$$

$$z^{[2](i)} = W^{[2]}a^{[1](i)} + b^{[2]} \quad (3)$$

$$\hat{y}^{(i)} = a^{[2](i)} = \sigma(z^{[2](i)}) \quad (4)$$

$$y_{prediction}^{(i)} = \begin{cases} 1 & \text{if } a^{[2](i)} > 0.5 \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

Given the predictions on all the examples, you can also compute the cost J as follows:

$$J = -\frac{1}{m} \sum_{i=0}^m (y^{(i)} \log(a^{[2](i)}) + (1 - y^{(i)}) \log(1 - a^{[2](i)})) \quad (6)$$

Reminder: The general methodology to build a Neural Network is to:

1. Define the neural network structure (# of input units, # of hidden units, etc).
2. Initialize the model's parameters
3. Loop:
 - Implement forward propagation
 - Compute loss
 - Implement backward propagation to get the gradients
 - Update parameters (gradient descent)

In practice, we'll often build helper functions to compute steps 1-3, test each function to make sure each one is working properly, and then merge them into one function called `nn_model()`. Once we've built `nn_model()` and learned the right parameters, we can make predictions on new data.

4.1 - Defining the neural network structure

Exercise 2 - layer_sizes

Define three variables: - `n_x`: the size of the input layer - `n_h`: the size of the hidden layer (**set this to 4, only for this Exercise 2**) - `n_y`: the size of the output layer

Hint: Use shapes of `X` and `Y` to find `n_x` and `n_y`. Also, hard code the hidden layer size to be 4.

```
[7]: # GRADED FUNCTION: layer_sizes
```

```
def layer_sizes(X, Y):  
    """
```

```

Arguments:
X -- input dataset of shape (input size, number of examples)
Y -- labels of shape (output size, number of examples)

Returns:
n_x -- the size of the input layer
n_h -- the size of the hidden layer
n_y -- the size of the output layer
"""

# YOUR CODE STARTS HERE


n_x = X.shape[0] # Number of features in the input layer
n_h = 4          # You can choose the size of the hidden layer (here,
↳ arbitrarily set to 4)
n_y = Y.shape[0] # Number of outputs

# YOUR CODE ENDS HERE
return (n_x, n_h, n_y)

```

```

[8]: t_X, t_Y = layer_sizes_test_case()
      (n_x, n_h, n_y) = layer_sizes(t_X, t_Y)
      print("The size of the input layer is: n_x = " + str(n_x))
      print("The size of the hidden layer is: n_h = " + str(n_h))
      print("The size of the output layer is: n_y = " + str(n_y))

      layer_sizes_test(layer_sizes)

```

```

The size of the input layer is: n_x = 5
The size of the hidden layer is: n_h = 4
The size of the output layer is: n_y = 2
All tests passed!

```

Expected output

```

The size of the input layer is: n_x = 5
The size of the hidden layer is: n_h = 4
The size of the output layer is: n_y = 2
All tests passed!
All tests passed.

```

4.2 - Initialize the model's parameters

Exercise 3 - initialize_parameters

Implement the function `initialize_parameters()`.

Instructions: - Make sure the parameters' sizes are right. Refer to the neural network figure above if needed. - Initialize the weights matrices with random values. - Use: `np.random.randn(a,b) * 0.01` to randomly initialize a matrix of shape (a,b). - Initialize the bias vectors as zeros. - Use: `np.zeros((a,b))` to initialize a matrix of shape (a,b) with zeros.

```
[9]: # GRADED FUNCTION: initialize_parameters

def initialize_parameters(n_x, n_h, n_y):
    """
    Argument:
    n_x -- size of the input layer
    n_h -- size of the hidden layer
    n_y -- size of the output layer

    Returns:
    params -- python dictionary containing your parameters:
                W1 -- weight matrix of shape (n_h, n_x)
                b1 -- bias vector of shape (n_h, 1)
                W2 -- weight matrix of shape (n_y, n_h)
                b2 -- bias vector of shape (n_y, 1)
    """

    # YOUR CODE STARTS HERE
    W1 = np.random.randn(n_h, n_x) * 0.01
    b1 = np.zeros((n_h, 1))
    W2 = np.random.randn(n_y, n_h) * 0.01
    b2 = np.zeros((n_y, 1))

    # YOUR CODE ENDS HERE

    parameters = {"W1": W1,
                  "b1": b1,
                  "W2": W2,
                  "b2": b2}

    return parameters
```

```
[10]: np.random.seed(2)
n_x, n_h, n_y = initialize_parameters_test_case()
parameters = initialize_parameters(n_x, n_h, n_y)

print("W1 = " + str(parameters["W1"]))
print("b1 = " + str(parameters["b1"]))
print("W2 = " + str(parameters["W2"]))
```



```
print("b2 = " + str(parameters["b2"]))

initialize_parameters_test(initialize_parameters)
```

```
W1 = [[-0.00416758 -0.00056267]
      [-0.02136196  0.01640271]
      [-0.01793436 -0.00841747]
      [ 0.00502881 -0.01245288]]
b1 = [[0.]
      [0.]
      [0.]
      [0.]]
W2 = [[-0.01057952 -0.00909008  0.00551454  0.02292208]]
b2 = [[0.]]
All tests passed!
```

Expected output

```
W1 = [[-0.00416758 -0.00056267]
      [-0.02136196  0.01640271]
      [-0.01793436 -0.00841747]
      [ 0.00502881 -0.01245288]]
b1 = [[0.]
      [0.]
      [0.]
      [0.]]
W2 = [[-0.01057952 -0.00909008  0.00551454  0.02292208]]
b2 = [[0.]]
All tests passed!
```

4.3 - The Loop

Exercise 4 - forward_propagation

Implement `forward_propagation()` using the following equations:

$$Z^{[1]} = W^{[1]}X + b^{[1]} \quad (1)$$

$$A^{[1]} = \tanh(Z^{[1]}) \quad (2)$$

$$Z^{[2]} = W^{[2]}A^{[1]} + b^{[2]} \quad (3)$$

$$\hat{Y} = A^{[2]} = \sigma(Z^{[2]}) \quad (4)$$

Instructions:

- Check the mathematical representation of the classifier in the figure above.
- Use the function `sigmoid()`. It's built into this notebook, (i.e., imported from the `planar_utils.py` module).
- Use the function `np.tanh()` from the numpy library.
- Implement using these steps:

1. Retrieve each parameter from the dictionary “parameters” (which is the output of `initialize_parameters()`) by using `parameters["."]`.
 2. Implement Forward Propagation. Compute $Z^{[1]}$, $A^{[1]}$, $Z^{[2]}$ and $A^{[2]}$ (the vector of all the predictions on all the examples in the training set).
- Values needed in the backpropagation are stored in “cache”. The cache will be given as an input to the backpropagation function.

```
[11]: # GRADED FUNCTION: forward_propagation

def forward_propagation(X, parameters):
    """
    Argument:
    X -- input data of size (n_x, m)
    parameters -- python dictionary containing your parameters (output of
    ↪ initialization function)

    Returns:
    A2 -- The sigmoid output of the second activation
    cache -- a dictionary containing "Z1", "A1", "Z2" and "A2"
    """
    # Retrieve each parameter from the dictionary "parameters". These are the
    ↪ W and b parameters from each
    # layer of our neural network that will allow forward prop from X to A2.

    # YOUR CODE STARTS HERE

    W1 = parameters["W1"]
    b1 = parameters["b1"]
    W2 = parameters["W2"]
    b2 = parameters["b2"]

    # YOUR CODE ENDS HERE

    # Implement Forward Propagation to calculate A2 (probabilities). These are
    ↪ the calculations of the nodes in
    # layers 1 and 2 of our neural network yielding our Z and A values for each
    ↪ of the layers, ending at A2.

    # YOUR CODE STARTS HERE

    # First layer computations
    Z1 = np.dot(W1, X) + b1          # Linear step for layer 1
    A1 = np.tanh(Z1)                 # Activation using tanh for layer 1
```

```

# Second layer computations
Z2 = np.dot(W2, A1) + b2      # Linear step for layer 2
A2 = sigmoid(Z2)              # Activation using sigmoid for output layer

# YOUR CODE ENDS HERE

assert(A2.shape == (1, X.shape[1]))

cache = {"Z1": Z1,
         "A1": A1,
         "Z2": Z2,
         "A2": A2}

return A2, cache

```

```

[12]: t_X, parameters = forward_propagation_test_case()
      A2, cache = forward_propagation(t_X, parameters)
      print("A2 = " + str(A2))

      forward_propagation_test(forward_propagation)

```

```

A2 = [[0.21292656 0.21274673 0.21295976]]
All tests passed!

```

Expected output

```

A2 = [[0.21292656 0.21274673 0.21295976]]
All tests passed!
All tests passed.

```

4.4 - Compute the Cost

Now that we've computed $A^{[2]}$ (in the Python variable "A2"), which contains $a^{[2](i)}$ for all examples, we can compute the cost function as follows:

$$J = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(a^{[2](i)}) + (1 - y^{(i)}) \log(1 - a^{[2](i)})) \quad (13)$$

Exercise 5 - compute_cost

Implement `compute_cost()` to compute the value of the cost J .

Instructions: - There are many ways to implement the loss. This is one way to implement one part of the equation without for loops: $-\sum_{i=1}^m y^{(i)} \log(a^{[2](i)})$:

```

logprobs = np.multiply(np.log(A2), Y)
cost = - np.sum(logprobs)

```

- Use that to build the whole expression of the cost function.

Notes:

- Use either 1. `np.multiply()` and then `np.sum()` or 2. `np.dot()`
- If you use `np.multiply` followed by `np.sum` the end result will be a type `float`, whereas if you use `np.dot`, the result will be a 2D numpy array.
- You can use `np.squeeze()` to remove redundant dimensions (in the case of single float, this will be reduced to a zero-dimension array).
- You can also cast the array as a type `float` using `float()`.

```
[13]: # GRADED FUNCTION: compute_cost

def compute_cost(A2, Y):
    """
    Computes the cross-entropy cost given in equation (13)

    Arguments:
    A2 -- The sigmoid output of the second activation, of shape (1, number of
    ↪examples)
    Y -- "true" labels vector of shape (1, number of examples)

    Returns:
    cost -- cross-entropy cost given equation (13)

    """

    m = Y.shape[1] # number of examples

    # Compute the cost
    # YOUR CODE STARTS HERE

    cost = -(1/m) * np.sum(Y * np.log(A2) + (1 - Y) * np.log(1 - A2))

    # YOUR CODE ENDS HERE

    cost = float(np.squeeze(cost)) # makes sure cost is the dimension we
    ↪expect.
                                     # E.g., turns [[17]] into 17

    return cost
```

```
[14]: A2, t_Y = compute_cost_test_case()
cost = compute_cost(A2, t_Y)
```

```
print("cost = " + str(compute_cost(A2, t_Y)))

compute_cost_test(compute_cost)
```

cost = 0.6930587610394646

All tests passed!

Expected output

cost = 0.6930587610394646

All tests passed!

All tests passed.

4.5 - Implement Backpropagation

Using the cache computed during forward propagation, we can now implement backward propagation.

Exercise 6 - backward_propagation

Implement the function `backward_propagation()`.

Instructions: Backpropagation is usually the hardest (most mathematical) part in deep learning. Below are the formulas for backpropagation. Use the six equations on the right side since we are building a vectorized implementation.

Figure 1: Backpropagation. Use the six equations on the right.

- Tips:
 - To compute dZ^1 we'll need to compute $g^{[1]'}(Z^{[1]})$. Since $g^{[1]}(.)$ is the tanh activation function, if $a = g^{[1]}(z)$ then $g^{[1]'}(z) = 1 - a^2$. So we can compute $g^{[1]'}(Z^{[1]})$ using `(1 - np.power(A1, 2))`.

```
[15]: # GRADED FUNCTION: backward_propagation

def backward_propagation(parameters, cache, X, Y):
    """
    Implement the backward propagation using the instructions above.

    Arguments:
    parameters -- python dictionary containing our parameters
    cache -- a dictionary containing "Z1", "A1", "Z2" and "A2".
    X -- input data of shape (2, number of examples)
    Y -- "true" labels vector of shape (1, number of examples)

    Returns:
    grads -- python dictionary containing your gradients with respect to
    ↪ different parameters
    """
    m = X.shape[1]

    # First, retrieve W1 and W2 from the dictionary "parameters".
```

```

# YOUR CODE STARTS HERE

W1 = parameters["W1"]
W2 = parameters["W2"]

# YOUR CODE ENDS HERE

# Next, retrieve A1 and A2 from dictionary "cache".
# YOUR CODE STARTS HERE

A1 = cache["A1"]
A2 = cache["A2"]

# YOUR CODE ENDS HERE

# Finally, implment backpropagation: calculate dW1, db1, dW2, db2.
# YOUR CODE STARTS HERE

dZ2 = A2 - Y #derivative for output
↳ layer
dW2 = (1/m) * np.dot(dZ2, A1.T) #derivative for W2
db2 = (1/m) * np.sum(dZ2, axis=1, keepdims=True) #derivative for b2

dA1 = np.dot(W2.T, dZ2) #backprop into hidden
↳ layer
dZ1 = dA1 * (1 - np.power(A1, 2)) #derivative for tanh
↳ activation (1-tanh^2)
dW1 = (1/m) * np.dot(dZ1, X.T) #derivative for W1
db1 = (1/m) * np.sum(dZ1, axis=1, keepdims=True) #derivative for b1

# YOUR CODE ENDS HERE

grads = {"dW1": dW1,
         "db1": db1,
         "dW2": dW2,
         "db2": db2}

return grads

```

```
[16]: parameters, cache, t_X, t_Y = backward_propagation_test_case()
```

```
grads = backward_propagation(parameters, cache, t_X, t_Y)
print ("dW1 = "+ str(grads["dW1"]))
print ("db1 = "+ str(grads["db1"]))
print ("dW2 = "+ str(grads["dW2"]))
print ("db2 = "+ str(grads["db2"]))

backward_propagation_test(backward_propagation)
```

```
dW1 = [[ 0.00301023 -0.00747267]
 [ 0.00257968 -0.00641288]
 [-0.00156892  0.003893  ]
 [-0.00652037  0.01618243]]
db1 = [[ 0.00176201]
 [ 0.00150995]
 [-0.00091736]
 [-0.00381422]]
dW2 = [[ 0.00078841  0.01765429 -0.00084166 -0.01022527]]
db2 = [[-0.16655712]]
All tests passed!
```

Expected output

```
dW1 = [[ 0.00301023 -0.00747267]
 [ 0.00257968 -0.00641288]
 [-0.00156892  0.003893  ]
 [-0.00652037  0.01618243]]
db1 = [[ 0.00176201]
 [ 0.00150995]
 [-0.00091736]
 [-0.00381422]]
dW2 = [[ 0.00078841  0.01765429 -0.00084166 -0.01022527]]
db2 = [[-0.16655712]]
All tests passed!
All tests passed.
```

4.6 - Update Parameters

Exercise 7 - update_parameters

Implement the update rule. Use gradient descent. We have to use (dW1, db1, dW2, db2) in order to update (W1, b1, W2, b2).

General gradient descent rule: $\theta = \theta - \alpha \frac{\partial J}{\partial \theta}$ where α is the learning rate and θ represents a parameter.

Figure 2: The gradient descent algorithm with a good learning rate (converging) and a bad learning rate (diverging). Images courtesy of Adam Harley.

Hint

- Use `copy.deepcopy(...)` when copying lists or dictionaries that are passed as parameters to functions. It avoids input parameters being modified within the function.

```
[17]: # GRADED FUNCTION: update_parameters

def update_parameters(parameters, grads, learning_rate = 1.2):
    """
    Updates parameters using the gradient descent update rule given above

    Arguments:
    parameters -- python dictionary containing your parameters
    grads -- python dictionary containing your gradients

    Returns:
    parameters -- python dictionary containing your updated parameters
    """
    # Retrieve a copy of each parameter from the dictionary "parameters". Use
    ↪ copy.deepcopy(...) for W1 and W2
    # YOUR CODE STARTS HERE

    W1 = copy.deepcopy(parameters["W1"])
    W2 = copy.deepcopy(parameters["W2"])
    b1 = copy.deepcopy(parameters["b1"])
    b2 = copy.deepcopy(parameters["b2"])

    # YOUR CODE ENDS HERE

    # Retrieve each gradient from the dictionary "grads"
    # YOUR CODE STARTS HERE

    dW1 = grads["dW1"]
    db1 = grads["db1"]
    dW2 = grads["dW2"]
    db2 = grads["db2"]

    # YOUR CODE ENDS HERE

    # Implement the gradient descent update for the parameters W and b
    ↪ parameter for each of the 2 layers.
    # YOUR CODE STARTS HERE

    W1 = W1 - learning_rate * dW1
    b1 = b1 - learning_rate * db1
```



```
W2 = W2 - learning_rate * dW2
b2 = b2 - learning_rate * db2
```

```
# YOUR CODE ENDS HERE
```

```
parameters = {"W1": W1,
              "b1": b1,
              "W2": W2,
              "b2": b2}
```

```
return parameters
```

```
[18]: parameters, grads = update_parameters_test_case()
parameters = update_parameters(parameters, grads)
```

```
print("W1 = " + str(parameters["W1"]))
print("b1 = " + str(parameters["b1"]))
print("W2 = " + str(parameters["W2"]))
print("b2 = " + str(parameters["b2"]))
```

```
update_parameters_test(update_parameters)
```

```
W1 = [[-0.00643025  0.01936718]
      [-0.02410458  0.03978052]
      [-0.01653973 -0.02096177]
      [ 0.01046864 -0.05990141]]
b1 = [[-1.02420756e-06]
      [ 1.27373948e-05]
      [ 8.32996807e-07]
      [-3.20136836e-06]]
W2 = [[-0.01041081 -0.04463285  0.01758031  0.04747113]]
b2 = [[0.00010457]]
All tests passed!
```

Expected output

```
W1 = [[-0.00643025  0.01936718]
      [-0.02410458  0.03978052]
      [-0.01653973 -0.02096177]
      [ 0.01046864 -0.05990141]]
b1 = [[-1.02420756e-06]
      [ 1.27373948e-05]
      [ 8.32996807e-07]
      [-3.20136836e-06]]
W2 = [[-0.01041081 -0.04463285  0.01758031  0.04747113]]
b2 = [[0.00010457]]
```

All tests passed!

All tests passed.

4.7 - Integration

Integrate the functions we've built above in `nn_model()`

Exercise 8 - `nn_model`

Build the neural network model in `nn_model()`.

Instructions: The neural network model has to use the previous functions in the right order.

```
[19]: # GRADED FUNCTION: nn_model

def nn_model(X, Y, n_h, num_iterations = 10000, print_cost=False):
    """
    Arguments:
    X -- dataset of shape (2, number of examples)
    Y -- labels of shape (1, number of examples)
    n_h -- size of the hidden layer
    num_iterations -- Number of iterations in gradient descent loop
    print_cost -- if True, print the cost every 1000 iterations

    Returns:
    parameters -- parameters learnt by the model. They can then be used to
    ↪ predict.
    """

    np.random.seed(3)
    n_x = layer_sizes(X, Y)[0]
    n_y = layer_sizes(X, Y)[2]

    # Initialize the parameters.
    # YOUR CODE STARTS HERE
    parameters = initialize_parameters(n_x, n_h, n_y)

    # YOUR CODE ENDS HERE

    # Loop (gradient descent)

    for i in range(0, num_iterations):

        # Implement the following steps by calling the functions we wrote above:
        ↪
```

```

# 1. Forward propagation. Inputs: "X, parameters". Outputs: "A2, cache".
↪
# 2. Cost calculation. Inputs: "A2, Y". Outputs: "cost".
# 3. Backpropagation. Inputs: "parameters, cache, X, Y". Outputs: ↪
↪ "grads".
# 4. Gradient descent parameter update. Inputs: "parameters, grads". ↪
↪ Outputs: "parameters".

# YOUR CODE STARTS HERE

A2, cache = forward_propagation(X, parameters)
cost = compute_cost(A2, Y)
grads = backward_propagation(parameters, cache, X, Y)
parameters = update_parameters(parameters, grads)

# YOUR CODE ENDS HERE

# Print the cost every 1000 iterations
if print_cost and i % 1000 == 0:
    print ("Cost after iteration %i: %f" %(i, cost))

return parameters

```

```
[20]: nn_model_test(nn_model)
```

```

Cost after iteration 0: 0.693086
Cost after iteration 1000: 0.000220
Cost after iteration 2000: 0.000108
Cost after iteration 3000: 0.000072
Cost after iteration 4000: 0.000054
Cost after iteration 5000: 0.000043
Cost after iteration 6000: 0.000036
Cost after iteration 7000: 0.000030
Cost after iteration 8000: 0.000027
Cost after iteration 9000: 0.000024
W1 = [[ 0.71392202  1.31281102]
      [-0.76411243 -1.41967065]
      [-0.75040545 -1.38857337]
      [ 0.56495575  1.04857776]]
b1 = [[-0.0073536 ]
      [ 0.01534663]
      [ 0.01262938]
      [ 0.00218135]]
W2 = [[ 2.82545815 -3.3063945 -3.16116615  1.8549574 ]]
b2 = [[0.00393452]]

```

All tests passed!

Expected output

```
Cost after iteration 0: 0.693198
Cost after iteration 1000: 0.000219
Cost after iteration 2000: 0.000108
...
Cost after iteration 8000: 0.000027
Cost after iteration 9000: 0.000024
W1 = [[ 0.56305445 -1.03925886]
      [ 0.7345426  -1.36286875]
      [-0.72533346  1.33753027]
      [ 0.74757629 -1.38274074]]
b1 = [[-0.22240654]
      [-0.34662093]
      [ 0.33663708]
      [-0.35296113]]
W2 = [[ 1.82196893  3.09657075 -2.98193564  3.19946508]]
b2 = [[0.21344644]]
All tests passed!
All tests passed.
```

5 - Test the Model

5.1 - Predict

Exercise 9 - predict

Predict with the model by building `predict()`. Use forward propagation to predict results.

Reminder: $y_{prediction} = \mathbb{1}_{activation > 0.5} = \begin{cases} 1 & \text{if } activation > 0.5 \\ 0 & \text{otherwise} \end{cases}$

As an example, if you would like to set the entries of a matrix `X` to 0 and 1 based on a threshold you would do: `X_new = (X > threshold)`

```
[21]: # GRADED FUNCTION: predict

def predict(parameters, X):
    """
    Using the learned parameters, predicts a class for each example in X

    Arguments:
    parameters -- python dictionary containing your parameters
    X -- input data of size (n_x, m)

    Returns
    predictions -- vector of predictions of our model (red: 0 / blue: 1)
    """
```

```
# Compute probabilities using forward propagation, and classifies to 0/1  
↪ using 0.5 as the threshold.
```

```
# YOUR CODE STARTS HERE  
#Compute probabilities using forward propagation  
A2, _ = forward_propagation(X, parameters)  
#Convert probabilities A2 to actual predictions using 0.5 as the threshold  
predictions = (A2 > 0.5).astype(int)  
  
# YOUR CODE ENDS HERE  
  
return predictions
```

```
[22]: parameters, t_X = predict_test_case()  
  
predictions = predict(parameters, t_X)  
print("Predictions: " + str(predictions))  
  
predict_test(predict)
```

```
Predictions: [[1 0 1]]  
All tests passed!
```

Expected output

```
Predictions: [[ True False  True]]  
All tests passed!  
All tests passed.
```

5.2 - Test the Model on the Planar Dataset


It's time to run the model and see how it performs on a planar dataset. Run the following code to test your model with a single hidden layer of n_h hidden units!

```
[23]: # Build a model with a n_h-dimensional hidden layer  
parameters = nn_model(X, Y, n_h = 4, num_iterations = 10000, print_cost=True)  
  
# Plot the decision boundary  
plot_decision_boundary(lambda x: predict(parameters, x.T), X, Y)  
plt.title("Decision Boundary for hidden layer size " + str(4))
```

```
Cost after iteration 0: 0.693162  
Cost after iteration 1000: 0.258625  
Cost after iteration 2000: 0.239334  
Cost after iteration 3000: 0.230802  
Cost after iteration 4000: 0.225528  
Cost after iteration 5000: 0.221845  
Cost after iteration 6000: 0.219094  
Cost after iteration 7000: 0.220668
```

```
Cost after iteration 8000: 0.219411
Cost after iteration 9000: 0.218486
```

```
[23]: Text(0.5, 1.0, 'Decision Boundary for hidden layer size 4')
```



```
[24]: # Print accuracy
predictions = predict(parameters, X)
print ('Accuracy: %d' % float((np.dot(Y, predictions.T) + np.dot(1 - Y, 1 -
↪ predictions.T)) / float(Y.size) * 100) + '%')
```

Accuracy: 90%

/tmp/ipykernel_13494/1304927518.py:3: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

```
print ('Accuracy: %d' % float((np.dot(Y, predictions.T) + np.dot(1 - Y, 1 -
predictions.T)) / float(Y.size) * 100) + '%')
```

Expected Output:

Accuracy

90%

Accuracy is really high compared to Logistic Regression. The model has learned the patterns of the flower's petals! Unlike logistic regression, neural networks are able to learn even highly non-linear decision boundaries.

1.1.1 Congrats - you're done!

Here's a quick recap of all you just accomplished:

- Built a complete 2-class classification neural network with a hidden layer
- Made good use of a non-linear unit
- Computed the loss
- Implemented forward and backward propagation

You've created a neural network that can learn patterns! Excellent work. Below is some code that shows some techniques for trying out different sizes for the hidden layer.

6 - Tuning hidden layer size (optional/ungraded exercise)

Run the following code (it may take 1-2 minutes). Then, observe different behaviors of the model for various hidden layer sizes.








```
[25]: # This may take about 2 minutes to run

plt.figure(figsize=(16, 32))
hidden_layer_sizes = [1, 2, 3, 4, 5, 20, 50]
for i, n_h in enumerate(hidden_layer_sizes):
    plt.subplot(5, 2, i+1)
    plt.title('Hidden Layer of size %d' % n_h)
    parameters = nn_model(X, Y, n_h, num_iterations = 5000)
    plot_decision_boundary(lambda x: predict(parameters, x.T), X, Y)
    predictions = predict(parameters, X)
    accuracy = float((np.dot(Y,predictions.T) + np.dot(1 - Y, 1 - predictions.
↪T)) / float(Y.size)*100)
    print ("Accuracy for {} hidden units: {}".format(n_h, accuracy))
```

```
/tmp/ipykernel_13494/164830050.py:11: DeprecationWarning: Conversion of an array
with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you
extract a single element from your array before performing this operation.
(Deprecated NumPy 1.25.)
```

```
    accuracy = float((np.dot(Y,predictions.T) + np.dot(1 - Y, 1 - predictions.T))
/ float(Y.size)*100)
```

```
Accuracy for 1 hidden units: 67.5 %
Accuracy for 2 hidden units: 67.25 %
Accuracy for 3 hidden units: 90.75 %
Accuracy for 4 hidden units: 90.5 %
Accuracy for 5 hidden units: 91.25 %
Accuracy for 20 hidden units: 90.75 %
Accuracy for 50 hidden units: 90.25 %
```



Interpretation: - The larger models (with more hidden units) are able to fit the training set better, until eventually the largest models overfit the data. - The best hidden layer size seems to be around $n_h = 5$. Indeed, a value around here seems to fit the data well without also incurring noticeable overfitting. - Later, we'll use regularization in neural networks which allows using large models (such as $n_h = 50$) without much overfitting.

References:

- <http://cs231n.github.io/neural-networks-case-study/>

[]: